

Hierarchical Heavy Hitter

TANG DING

Papers

- Mitzenmacher, M., Steinke, T., & Thaler, J. Hierarchical heavy hitters with the space saving algorithm. (*ALENEX 2012*)
- Ben Basat, R., Einziger, G., Friedman, R., Luizelli, M. C., & Waisbard, E. Constant time updates in hierarchical heavy hitters. (*SIGCOMM 2017*)

Application


- Network traffic monitoring
- Anomaly detection
- DDoS detection.

Heavy Hitter


- Basically, equals to frequent items
- Identifies frequent:
 - Source.
 - Source-Destination pairs.

Example

We want to find Source IP whose frequency ≥ 2



Hierarchical Heavy Hitter



Source

$220.7.16.*$




Destination

1D Hierarchy example


- Level4: *.*.*.*
- Level3: 21.*.*.*
- Level2: 21.132.*.*
- Level1: 21.132.145.*, 21.132.146.*
- Level0: 21.132.145.13, 21.132.146.14

2D Hierarchy example



Example

We want to find HHH whose frequency ≥ 2




Why count of 21.132.161.* is 2?

Problems

- Often, the large volume of network traffic makes it infeasible to store the relevant data in memory.
- Space-saving streaming algorithm.

Example

We want to find HHH whose frequency ≥ 2



Output confidence interval instead of actual frequency

Target

- Accuracy: Limited Errors
- Coverage: No False Negatives

Target

- Given parameter ϵ , we want to output a set of items S , and lower and upper bounds of frequency ℓ and u , such that they satisfy two properties, as follows:
 - 1. Accuracy. ℓ and u for all i .
 - 2. Coverage. For set S , ℓ . Define the conditioned count of i with respect to S to be c_i . We require for all prefixes

Algorithm

- Two main procedures:
 - Update
 - Output

Update

Time Step	Update	Counter 1	Counter 2	Counter 3
0		<i>unused</i>	unused	unused
1	(a,+2)	(a,2)	<i>unused</i>	unused
2	(b,+6)	(a,2)	(b,6)	<i>unused</i>
3	(c,+4)	(<i>a</i> ,2)	(b,6)	(c,4)
4	(a,+3)	(a,5)	(b,6)	(<i>c</i> ,4)
5	(d,+4)	(<i>a</i> ,5)	(b,6)	(d,8)
6	(e,+4)	(e,9)	(<i>b</i> ,6)	(d,8)

For any item C being tracked, the actual frequency $f(C)$ is in $(N_c - N_{\min}, N_c)$,

where N_c is the count of that item, and N_{\min} is the minimum count in all counters

Update

Time T	Counter 1	Counter 2	Counter 3
Level 1	(a.*.*., 20)	(h.*.*., 12)	(q.*.*., 16)
Level 2	(a.b.*.*., 18)	(h.i.*.*., 12)	(q.r.*.*., 16)
Level 3	(a.b.c.*., 18)	(q.r.s.*., 9)	(h.i.j.*., 14)
Level 4	(a.b.c.d, 10)	(h.i.m.n, 15)	(a.b.c.e, 8)

(w.x.y.z, +3)


Time T+1	Counter 1	Counter 2	Counter 3
Level 1	(a.*.*., 20)	(w.*.*., 15)	(q.*.*., 16)
Level 2	(a.b.*.*., 18)	(w.x.*.*., 15)	(q.r.*.*., 16)
Level 3	(a.b.c.*., 18)	(w.x.y.*., 12)	(h.i.j.*., 14)
Level 4	(a.b.c.d, 10)	(h.i.m.n, 15)	(w.x.y.z, 11)

Update

- For each level, set counters
- Three procedures
 - Find counter for certain IP (pairs)
 - Find counter with minimum count
 - Update counter found (item / count)
- Update Complexity
 - Arbitrary Update: Using heaps –
 - Unitary Update:
 - is the number of levels

Unitary Update

- Data Structure
 - HashTable
 - Item{HashEntry, GroupEntry, NextItemInGroup, PreviousItemInGroup}
 - Group {Items, Count, PreviousGroup, NextGroup}
 - Sorted link list with entries of same count combined



(a) Stream-Summary,
 $S = X, Y$

(b) Stream-Summary,
 $S = X, Y, Y$

(c) Stream-Summary,
 $S = X, Y, Y, Z$

Output (for 1D)

OUTPUTHH1D(threshold ϕ)

```
1  /* par( $e$ ) is parent of  $e^*$  /
2  Let  $s_e = 0$  for all  $e$ 
3  /*  $s_e$  conservatively estimates the difference
   between unconditioned and conditioned counts of  $e^*$  /
4  for each  $e$  in postorder
5      ( $f_{\min}(e), f_{\max}(e)$ ) = GetEstimateSS( $SS(n), e$ )
6      if  $f_{\max}(e) - s_e \geq \phi N$ 
7          print( $e, f_{\min}(e), f_{\max}(e)$ )
8           $s_{\text{par}(e)} + = f_{\min}(e)$ 
9      else  $s_{\text{par}(e)} + = s_e$ 
```

Proof


- By using $\log n$ space ($\log n$ counters in each level), we can achieve our target in both accuracy and coverage
 - Accuracy:
 - if $\epsilon > 0$,
 - $\epsilon = \frac{1}{n}$
 - Coverage

Output (for 2D)

OUTPUTHHH2D(threshold ϕ)


```
1   $P = \emptyset$ 
2  for level  $l=L$  downto 0
3      for each item  $p$  at level  $l$ 
4          Let  $n$  be the lattice node that  $p$  belongs to
5           $(f_{\min}(p), f_{\max}(p)) = \text{GetEstimateSS}(SS(n), p)$ 
6           $F'_p = f_{\max}(p)$ 
7           $H_p = \{h \in P \text{ such that } \nexists h' \in P : h \prec h' \prec p\}$ 
8          for each  $h \in H_p$ 
9               $F'_p = F'_p - f_{\min}(h)$ 
10         for each pair of distinct elements  $h, h'$  in  $H_p$ 
11              $q = \text{glb}(h, h')$ 
12             if  $\nexists h_3 \neq h, h'$  in  $H_p$  s.t.  $q \preceq h_3$ 
13                  $F'_p = F'_p + f_{\max}(q)$ 
14         if  $F'_p \geq \phi N$ 
15              $P = P \cup \{p\}$ 
16         print( $p, f_{\min}(p), f_{\max}(p)$ )
```

Output



Randomized

Select a prefix at random and count it



Additional Speedup

188.3.12.3

181.7.20.3

92.67.7.81

181.7.20.6

188.67.7.1

181.7.20.2

With
some
probability

Compute a
random
prefix

181.7.20.*



Ignore
packet

Level0 Counting

Level1 Counting

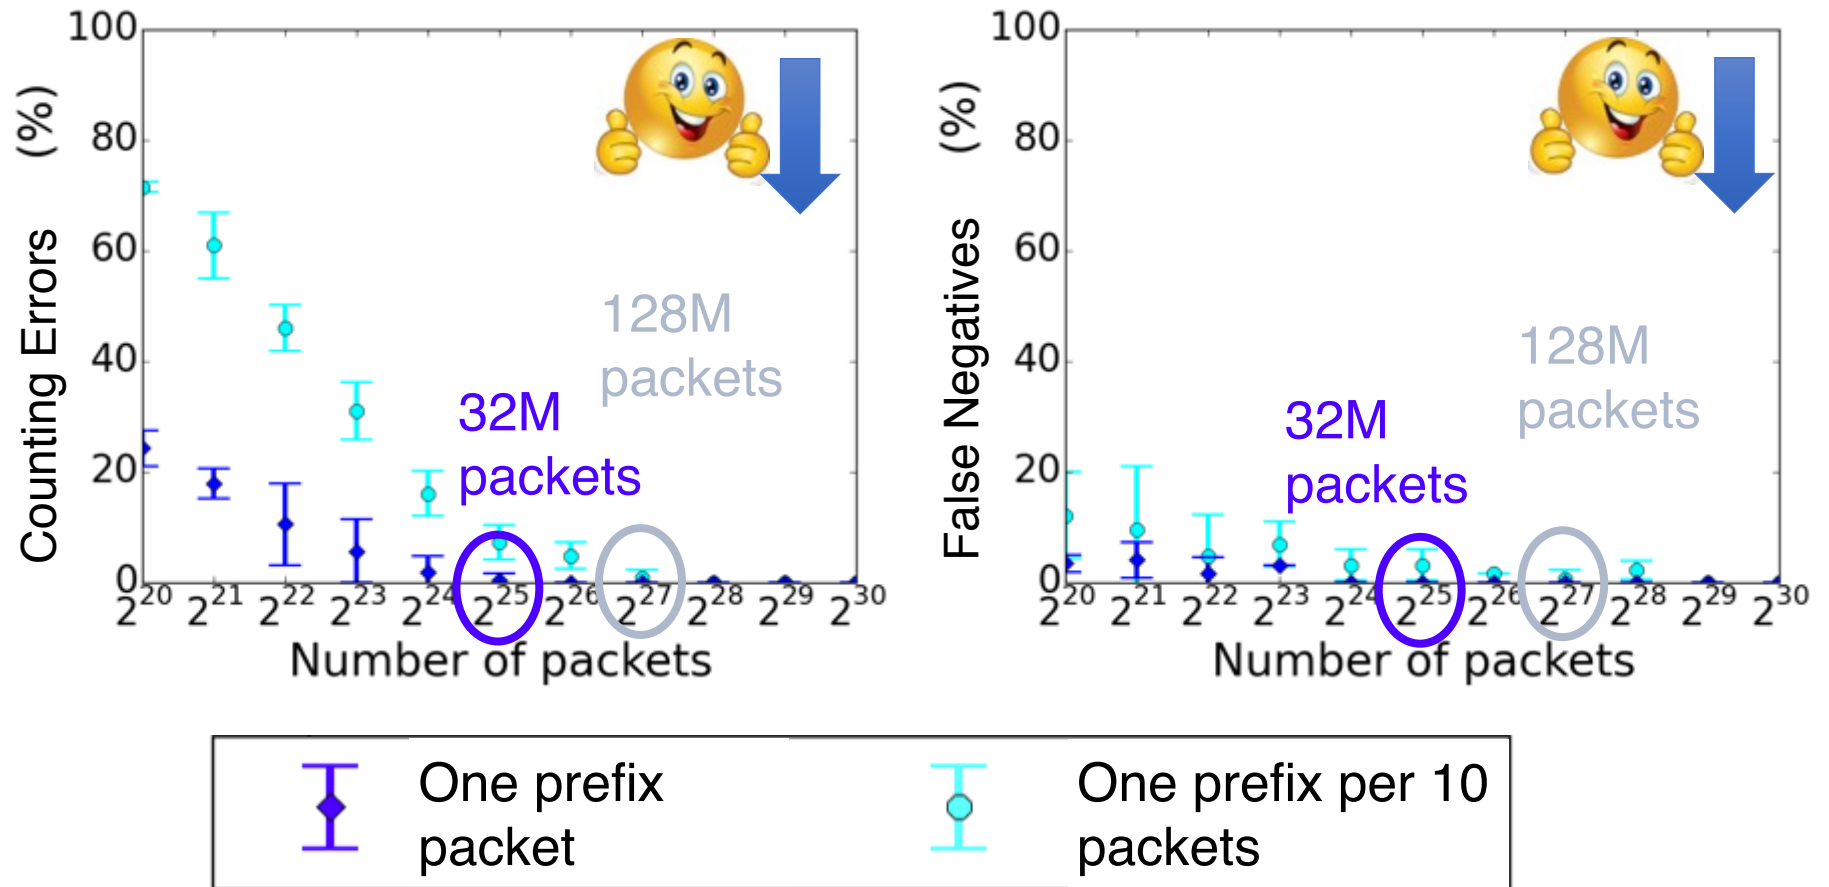
Level2 Counting

Level3 Counting


Level4 Counting

Accuracy

“Accuracy improves with the number of packets”



Performance



One update per packet



One update per 10 packets



Partial Ancestry

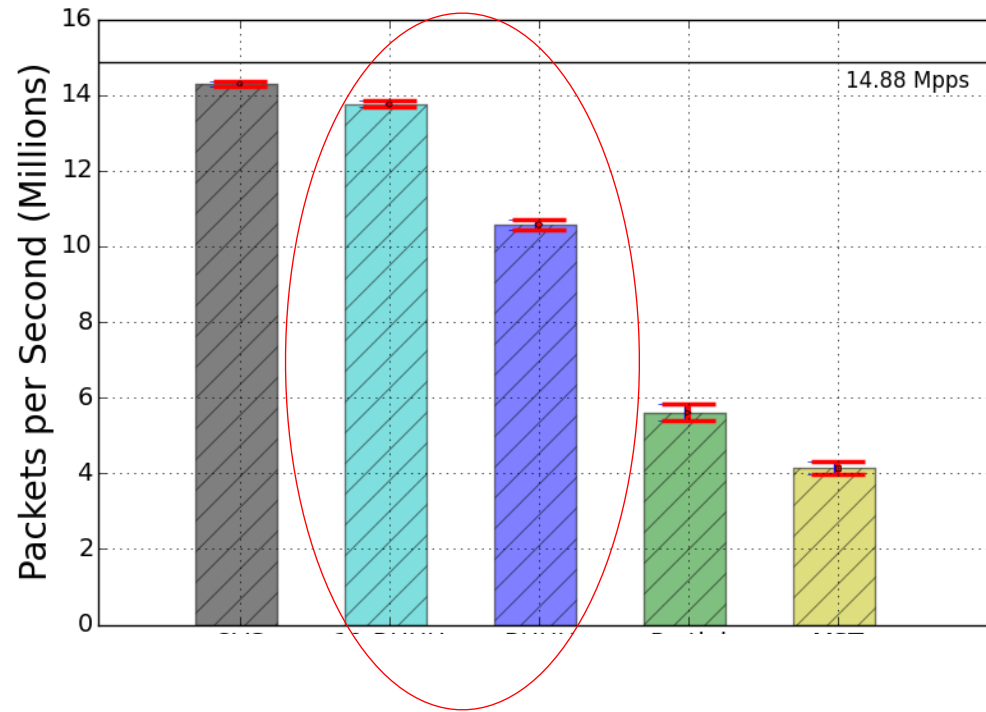


Full Ancestry



Mitzenmacher et al.

Performance



Highlights:

- Only -4% overheads for HHH in the Open vSwitch data plane!
- +250% throughput improvement compared to previous work.

